

✓ Install Dependencies

(Remember to choose GPU in Runtime if not already selected. Runtime --> Change Runtime Type --> Hardware accelerator --> GPU)

```
# clone YOLOv5 repository
!git clone https://github.com/ultralytics/yolov5.git # clone repo
%cd yolov5

Cloning into 'yolov5'...
remote: Enumerating objects: 16413, done.
remote: Counting objects: 100% (5/5), done.
remote: Compressing objects: 100% (5/5), done.
remote: Total 16413 (delta 0), reused 4 (delta 0), pack-reused 16408
Receiving objects: 100% (16413/16413), 14.96 MiB | 10.23 MiB/s, done.
Resolving deltas: 100% (11263/11263), done.
/content/yolov5

# install dependencies as necessary
!pip install -qr requirements.txt # install dependencies (ignore errors)
import torch

from IPython.display import Image, clear_output # to display images
# from utils.google_utils import gdrive_download # to download models/datasets

# clear_output()
print('Setup complete. Using torch %s %s' % (torch.__version__, torch.cuda.get_device_properties(0) if torch.cuda.is_available() else None))

_____  
196.4/196.4 kB 5.3 MB/s eta 0:00:00  
_____  
709.5/709.5 kB 14.2 MB/s eta 0:00:00  
_____  
62.7/62.7 kB 8.1 MB/s eta 0:00:00

Setup complete. Using torch 2.1.0+cu121 _CudaDeviceProperties(name='Tesla T4', major=7, minor=5, total_memory=15102MB, multi_pr
```

✓ Your Custom Data

```
# Export code snippet and paste here
%cd /content
!wget https://github.com/aadityaadyot/CV-Hand-Signal-Project/raw/main/Data/Sign_language_data.zip
!unzip -q Sign_language_data.zip

/content
--2024-02-09 20:33:29-- https://github.com/aadityaadyot/CV-Hand-Signal-Project/raw/main/Data/Sign\_language\_data.zip
Resolving github.com (github.com)... 192.30.255.112
Connecting to github.com (github.com)|192.30.255.112|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://raw.githubusercontent.com/aadityaadyot/CV-Hand-Signal-Project/main/Data/Sign\_language\_data.zip [following]
--2024-02-09 20:33:29-- https://raw.githubusercontent.com/aadityaadyot/CV-Hand-Signal-Project/main/Data/Sign\_language\_data.zip
Resolving raw.githubusercontent.com (raw.githubusercontent.com)... 185.199.108.133, 185.199.110.133, 185.199.111.133, ...
Connecting to raw.githubusercontent.com (raw.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 39777787 (38M) [application/zip]
Saving to: 'Sign_language_data.zip'

Sign_language_data. 100%[=====] 37.93M --.-KB/s   in 0.1s

2024-02-09 20:33:30 (303 MB/s) - 'Sign_language_data.zip' saved [39777787/39777787]
```



```
# this is the YAML file we're loading into this notebook with our data
%cat data.yaml

train: ../train/images
val: ../test/images

nc: 7
names: ['Fuck Off', 'Hello', 'ILoveYou', 'Namaste', 'Please', 'Thanks', 'Yes']
```

✓ Define Model Configuration and Architecture

We will write a yaml script that defines the parameters for our model like the number of classes, anchors, and each layer.

You do not need to edit these cells, but you may.

```
# define number of classes based on YAML
import yaml
with open("data.yaml", 'r') as stream:
    num_classes = str(yaml.safe_load(stream)['nc'])

num_classes
'7'

#this is the model configuration we will use
%cat /content/yolov5/models/yolov5s.yaml

# YOLOv5 🚀 by Ultralytics, AGPL-3.0 license

# Parameters
nc: 80 # number of classes
depth_multiple: 0.33 # model depth multiple
width_multiple: 0.50 # layer channel multiple
anchors:
- [10, 13, 16, 30, 33, 23] # P3/8
- [30, 61, 62, 45, 59, 119] # P4/16
- [116, 90, 156, 198, 373, 326] # P5/32

# YOLOv5 v6.0 backbone
backbone:
    # [from, number, module, args]
    [
        [-1, 1, Conv, [64, 6, 2, 2]], # 0-P1/2
        [-1, 1, Conv, [128, 3, 2]], # 1-P2/4
        [-1, 3, C3, [128]],
    ]
```

```

[-1, 1, Conv, [256, 3, 2]], # 3-P3/8
[-1, 6, C3, [256]],
[-1, 1, Conv, [512, 3, 2]], # 5-P4/16
[-1, 9, C3, [512]],
[-1, 1, Conv, [1024, 3, 2]], # 7-P5/32
[-1, 3, C3, [1024]],
[-1, 1, SPPF, [1024, 5]], # 9
]

# YOLOv5 v6.0 head
head: [
    [-1, 1, Conv, [512, 1, 1]],
    [-1, 1, nn.Upsample, [None, 2, "nearest"]],
    [[[-1, 6], 1, Concat, [1]]], # cat backbone P4
    [-1, 3, C3, [512, False]], # 13

    [-1, 1, Conv, [256, 1, 1]],
    [-1, 1, nn.Upsample, [None, 2, "nearest"]],
    [[[-1, 4], 1, Concat, [1]]], # cat backbone P3
    [-1, 3, C3, [256, False]], # 17 (P3/8-small)

    [-1, 1, Conv, [256, 3, 2]],
    [[[-1, 14], 1, Concat, [1]]], # cat head P4
    [-1, 3, C3, [512, False]], # 20 (P4/16-medium)

    [-1, 1, Conv, [512, 3, 2]],
    [[[-1, 10], 1, Concat, [1]]], # cat head P5
    [-1, 3, C3, [1024, False]], # 23 (P5/32-large)

    [[17, 20, 23], 1, Detect, [nc, anchors]], # Detect(P3, P4, P5)
]
]

#customize iPython writefile so we can write variables
from IPython.core.magic import register_line_cell_magic

@register_line_cell_magic
def writetemplate(line, cell):
    with open(line, 'w') as f:
        f.write(cell.format(**globals()))

```



```
%%writetemplate /content/yolov5/models/custom_yolov5s.yaml

# parameters
nc: {num_classes} # number of classes
depth_multiple: 0.33 # model depth multiple
width_multiple: 0.50 # layer channel multiple

# anchors
anchors:
- [10,13, 16,30, 33,23] # P3/8
- [30,61, 62,45, 59,119] # P4/16
- [116,90, 156,198, 373,326] # P5/32

# YOLOv5 backbone
backbone:
# [from, number, module, args]
[[[-1, 1, Focus, [64, 3]], # 0-P1/2
 [-1, 1, Conv, [128, 3, 2]], # 1-P2/4
 [-1, 3, BottleneckCSP, [128]],
 [-1, 1, Conv, [256, 3, 2]], # 3-P3/8
 [-1, 9, BottleneckCSP, [256]],
 [-1, 1, Conv, [512, 3, 2]], # 5-P4/16
 [-1, 9, BottleneckCSP, [512]],
 [-1, 1, Conv, [1024, 3, 2]], # 7-P5/32
 [-1, 1, SPP, [1024, [5, 9, 13]]],
 [-1, 3, BottleneckCSP, [1024, False]], # 9
]

# YOLOv5 head
head:
[[[-1, 1, Conv, [512, 1, 1]],
 [-1, 1, nn.Upsample, [None, 2, 'nearest']],
 [[-1, 6], 1, Concat, [1]], # cat backbone P4
 [-1, 3, BottleneckCSP, [512, False]], # 13

 [-1, 1, Conv, [256, 1, 1]],
 [-1, 1, nn.Upsample, [None, 2, 'nearest']]
```

```
[[[-1, 4], 1, Concat, [1]], # cat backbone P3
 [-1, 3, BottleneckCSP, [256, False]], # 17 (P3/8-small)

 [-1, 1, Conv, [256, 3, 2]],
 [[-1, 14], 1, Concat, [1]], # cat head P4
 [-1, 3, BottleneckCSP, [512, False]], # 20 (P4/16-medium)

 [-1, 1, Conv, [512, 3, 2]],
 [[-1, 10], 1, Concat, [1]], # cat head P5
 [-1, 3, BottleneckCSP, [1024, False]], # 23 (P5/32-large)

 [[17, 20, 23], 1, Detect, [nc, anchors]], # Detect(P3, P4, P5)
 ]
```

▼ Train Custom YOLOv5 Detector

Next, we'll fire off training!

Here, we are able to pass a number of arguments:

- **img**: define input image size
- **batch**: determine batch size
- **epochs**: define the number of training epochs. (Note: often, 3000+ are common here!)
- **data**: set the path to our yaml file
- **cfg**: specify our model configuration
- **weights**: specify a custom path to weights. (Note: you can download weights from the Ultralytics Google Drive [folder](#))
- **name**: result names
- **nosave**: only save the final checkpoint
- **cache**: cache images for faster training

```
# train yolov5s on custom data for 100 epochs
```

```
# time its performance
%time
%cd /content/yolov5/
!python train.py --img 545 --batch 16 --epochs 500 --data '../data.yaml' --cfg ./models/custom_yolov5s.yaml --weights yolov5s.pt
```

| | | | | | | |
|-----|----|----|------|---|-------|-------|
| all | 37 | 37 | 0.97 | 1 | 0.977 | 0.832 |
|-----|----|----|------|---|-------|-------|

500 epochs completed in 0.438 hours.
 Optimizer stripped from runs/train/yolov5s_results/weights/last.pt, 14.8MB
 Optimizer stripped from runs/train/yolov5s_results/weights/best.pt, 14.8MB

Validating runs/train/yolov5s_results/weights/best.pt...
 Fusing layers...
 custom_YOLOv5s summary: 182 layers, 7262700 parameters, 0 gradients

| Class | Images | Instances | P | R | mAP50 | mAP50-95: 100% 2/2 [00:00<00:00, 3.25it/s] |
|----------|--------|-----------|-------|---|-------|--|
| all | 37 | 37 | 0.968 | 1 | 0.977 | 0.85 |
| Fuck Off | 37 | 4 | 0.984 | 1 | 0.995 | 0.937 |
| Hello | 37 | 5 | 0.984 | 1 | 0.995 | 0.879 |
| ILoveYou | 37 | 7 | 0.865 | 1 | 0.871 | 0.732 |
| Namaste | 37 | 6 | 0.993 | 1 | 0.995 | 0.736 |
| Please | 37 | 4 | 0.985 | 1 | 0.995 | 0.92 |
| Thanks | 37 | 6 | 0.986 | 1 | 0.995 | 0.885 |
| Yes | 37 | 5 | 0.982 | 1 | 0.995 | 0.862 |

Results saved to runs/train/yolov5s_results
 CPU times: user 16.8 s, sys: 1.98 s, total: 18.8 s
 Wall time: 26min 52s

▼ Evaluate Custom YOLOv5 Detector Performance

Double-click (or enter) to edit

Training losses and performance metrics are saved to Tensorboard and also to a logfile defined above with the **--name** flag when we train. In our case, we named this `yolov5s_results`. (If given no name, it defaults to `results.txt`.) The results file is plotted as a png after training completes.

Note from Glenn: Partially completed `results.txt` files can be plotted with `from utils.utils import plot_results; plot_results()`.

```
# Start tensorboard
# Launch after you have started training
# logs save in the folder "runs"
%load_ext tensorboard
%tensorboard --logdir runs
```

Filter runs (regex) Filter tags (regex)

All

Scalars

Image

Histogram

 Settings Run ↑ Pinned train/yolov5s_results

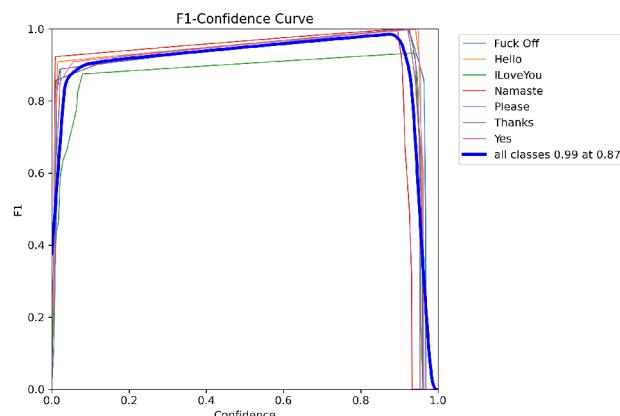
Pin cards for a quick view and comparison

F1_curve

F1_curve

train/yolov5s_resul...

Step 499



PR_curve

PR_curve

train/yolov5s_resul...

Step 499

Settings

GENERAL

Horizontal Axis

Step

 Enable step selection and data table
(Scalars only) Enable Range Selection Link by step 499

Card Width

 100%

SCALARS

Smoothing

 0.6

Tooltip sorting method

Alphabetical

 Ignore outliers in chart scaling Partition non-monotonic X axis



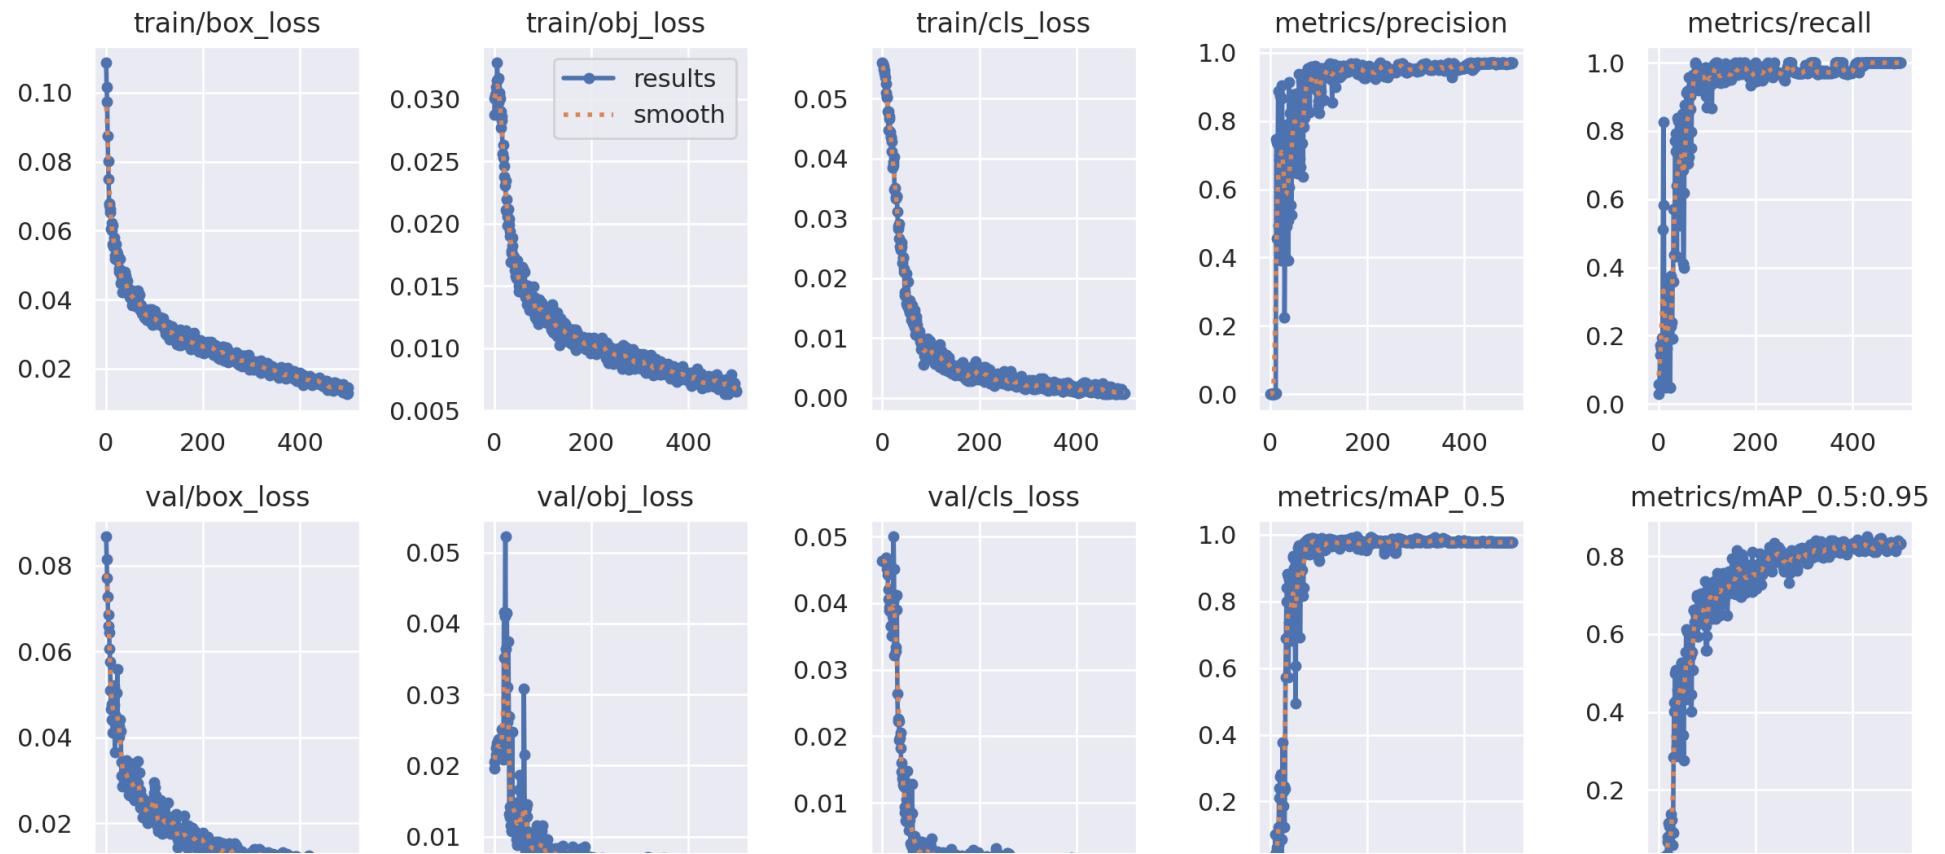
```
!nvidia-smi
```

```
Fri Feb  9 21:11:57 2024
```

| NVIDIA-SMI 535.104.05 | | | Driver Version: 535.104.05 | | CUDA Version: 12.2 | | |
|-----------------------|----------|---------------|----------------------------|-----------------|--------------------|------------|--------|
| GPU | Name | Persistence-M | Bus-Id | Disp.A | Volatile | Uncorr. | ECC |
| Fan | Temp | Perf | Pwr:Usage/Cap | Memory-Usage | GPU-Util | Compute M. | MIG M. |
| 0 | Tesla T4 | Off | 00000000:00:04.0 | Off | | 0 | |
| N/A | 49C | P8 | 10W / 70W | 3MiB / 15360MiB | 0% | Default | N/A |

| Processes: | | | | | |
|----------------------------|----|----|-----|------|------------------|
| GPU | GI | CI | PID | Type | Process name |
| ID | | | ID | ID | GPU Memory Usage |
| No running processes found | | | | | |

```
# we can also output some older school graphs if the tensor board isn't working for whatever reason...
from utils.plots import plot_results # plot results.txt as results.png
Image(filename='/content/yolov5/runs/train/yolov5s_results/results.png', width=1000) # view results.png
```



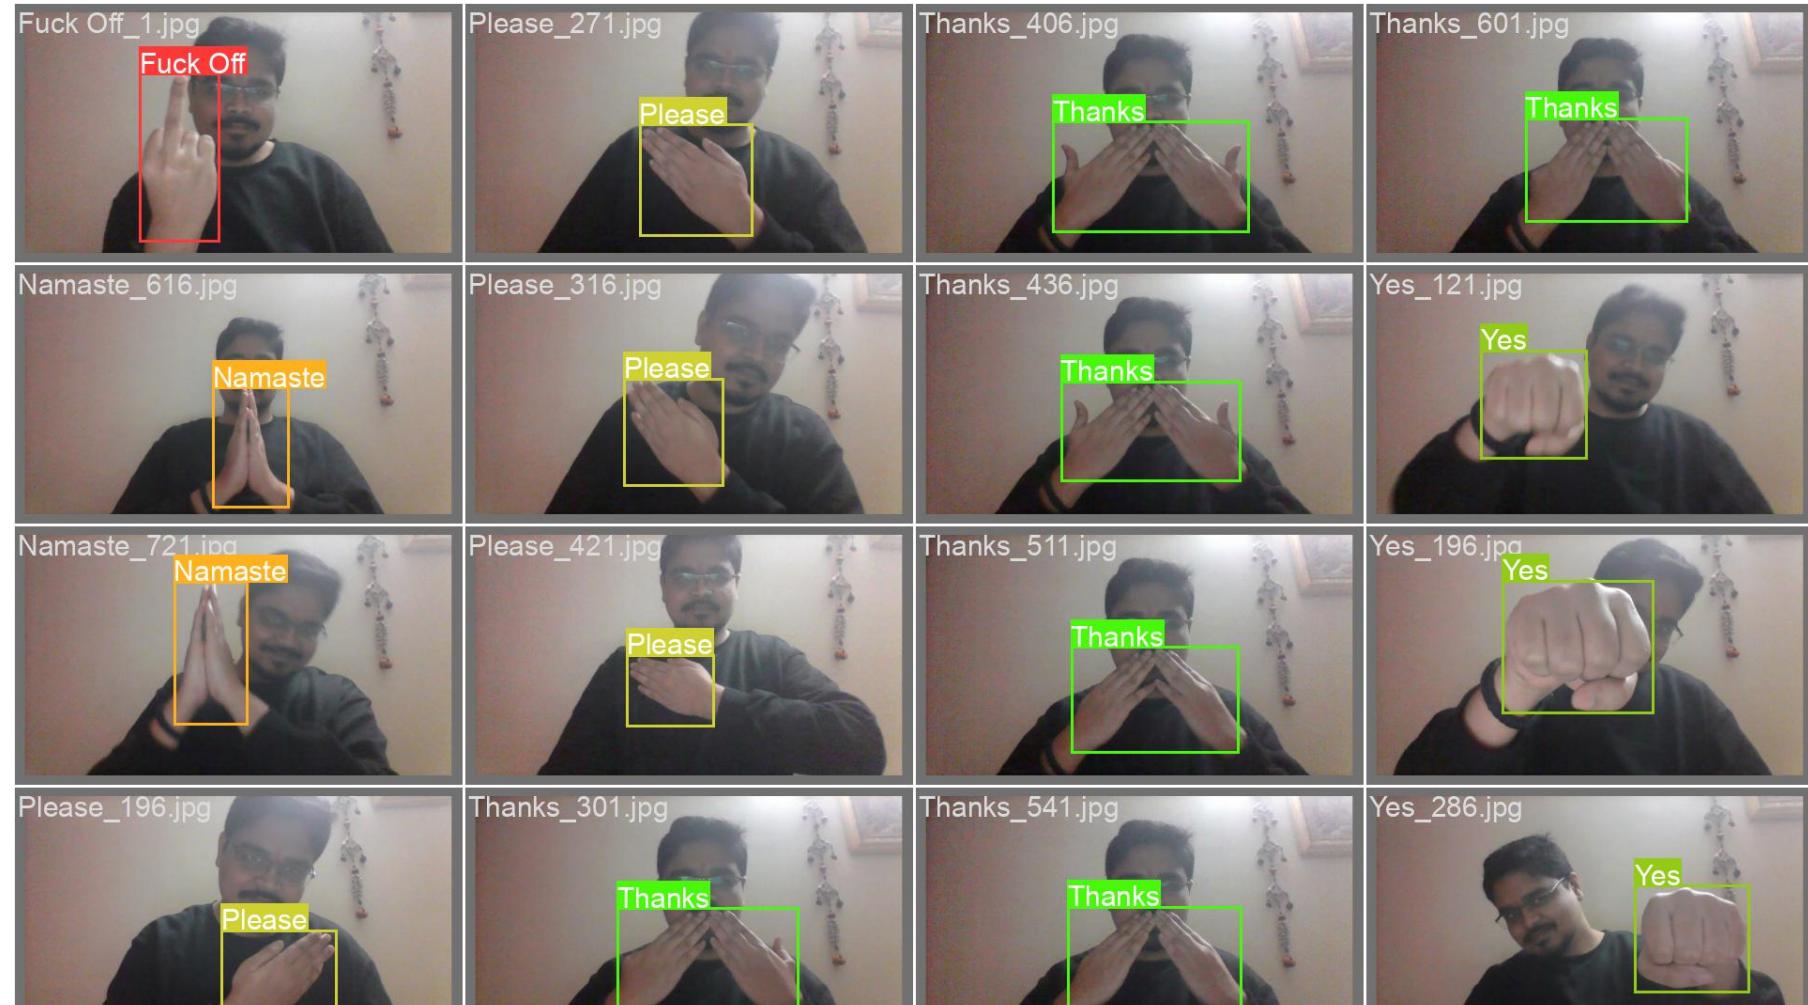
▼ Curious? Visualize Our Training Data with Labels

After training starts, view `train*.jpg` images to see training images, labels and augmentation effects.

Note a mosaic dataloader is used for training (shown below), a new dataloading concept developed by Glenn Jocher and first featured in [YOLOv4](#).

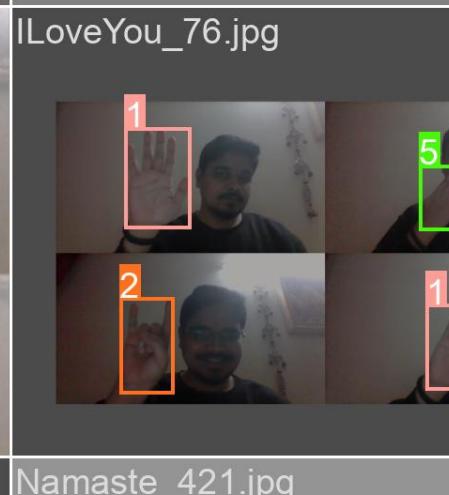
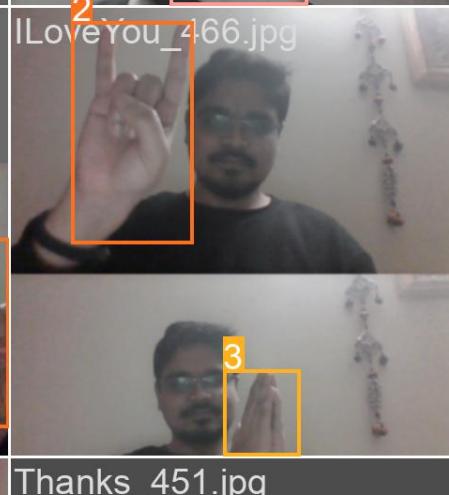
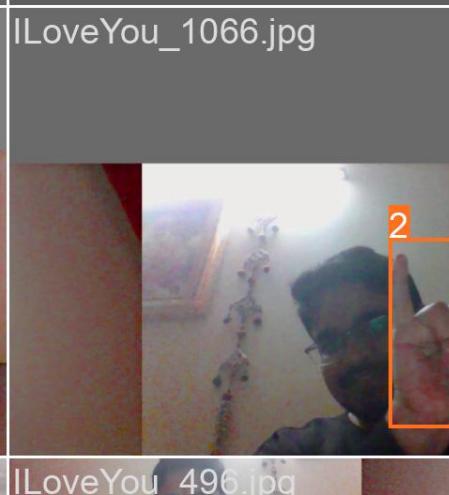
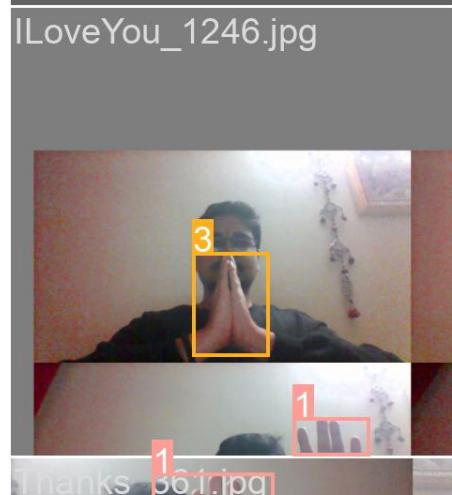
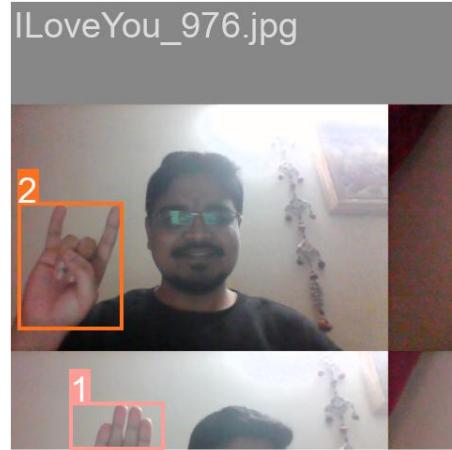
```
# first, display our ground truth data
print("GROUND TRUTH TRAINING DATA:")
Image(filename='/content/yolov5/runs/train/yolov5s_results/val_batch0_labels.jpg', width=900)
```

GROUND TRUTH TRAINING DATA:



```
# print out an augmented training example
print("GROUND TRUTH AUGMENTED TRAINING DATA:")
Image(filename='/content/yolov5/runs/train/yolov5s_results/train_batch0.jpg', width=900)
```

GROUND TRUTH AUGMENTED TRAINING DATA:

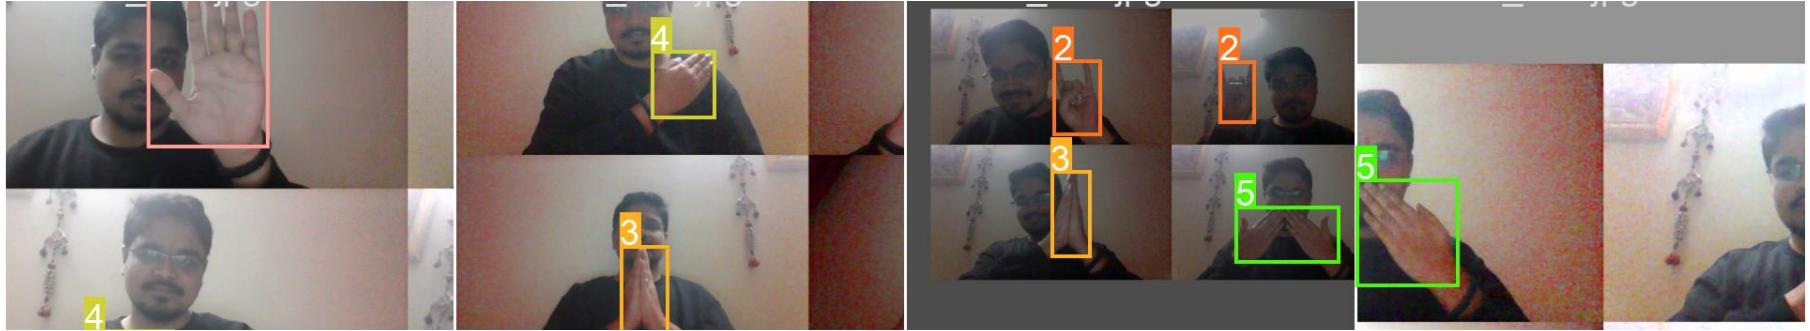


Thanks_361.jpg

ILoveYou_496.jpg

Thanks_451.jpg

Namaste_421.jpg



✓ Run Inference With Trained Weights

Run inference with a pretrained checkpoint on contents of `test/images` folder downloaded from Roboflow.

```
# trained weights are saved by default in our weights folder  
%ls runs/
```

```
train/
```

```
%ls runs/train/yolov5s_results/weights  
  
best.pt last.pt
```

Double-click (or enter) to edit

```
# when we ran this, we saw .007 second inference time. That is 140 FPS on a TESLA P100!  
# use the best weights!  
%cd /content/yolov5/  
!python detect.py --weights runs/train/yolov5s_results/weights/best.pt --img 545 --conf 0.5 --source ..  
/content/yolov5  
detect: weights=['runs/train/yolov5s_results/weights/best.pt'], source=..//test/images, data=data/coco128.yaml, imgsz=[545, 545]  
YOLOv5 🚀 v7.0-284-g95ebf68f Python-3.10.12 torch-2.1.0+cu121 CUDA:0 (Tesla T4, 15102MiB)
```

Fusing layers...

```
custom_YOLOv5s summary: 182 layers, 7262700 parameters, 0 gradients
WARNING ⚠ --img-size [545, 545] must be multiple of max stride 32, updating to [576, 576]
WARNING ⚠ NMS time limit 0.550s exceeded
image 1/37 /content/test/images/Fuck Off_1.jpg: 352x576 1 Fuck Off, 47.7ms
image 2/37 /content/test/images/Fuck Off_16.jpg: 352x576 1 Fuck Off, 9.0ms
image 3/37 /content/test/images/Fuck Off_31.jpg: 352x576 1 Fuck Off, 7.9ms
image 4/37 /content/test/images/Fuck Off_46.jpg: 352x576 1 Fuck Off, 7.9ms
image 5/37 /content/test/images>Hello_121.jpg: 352x576 1 Hello, 7.9ms
image 6/37 /content/test/images>Hello_161.jpg: 352x576 1 Hello, 7.8ms
image 7/37 /content/test/images>Hello_21.jpg: 352x576 1 Hello, 7.8ms
image 8/37 /content/test/images>Hello_31.jpg: 352x576 1 Hello, 7.9ms
image 9/37 /content/test/images>Hello_61.jpg: 352x576 1 Hello, 8.8ms
image 10/37 /content/test/images/I Love You_136.jpg: 352x576 1 I Love You, 7.9ms
image 11/37 /content/test/images/I Love You_16.jpg: 352x576 1 I Love You, 7.9ms
```